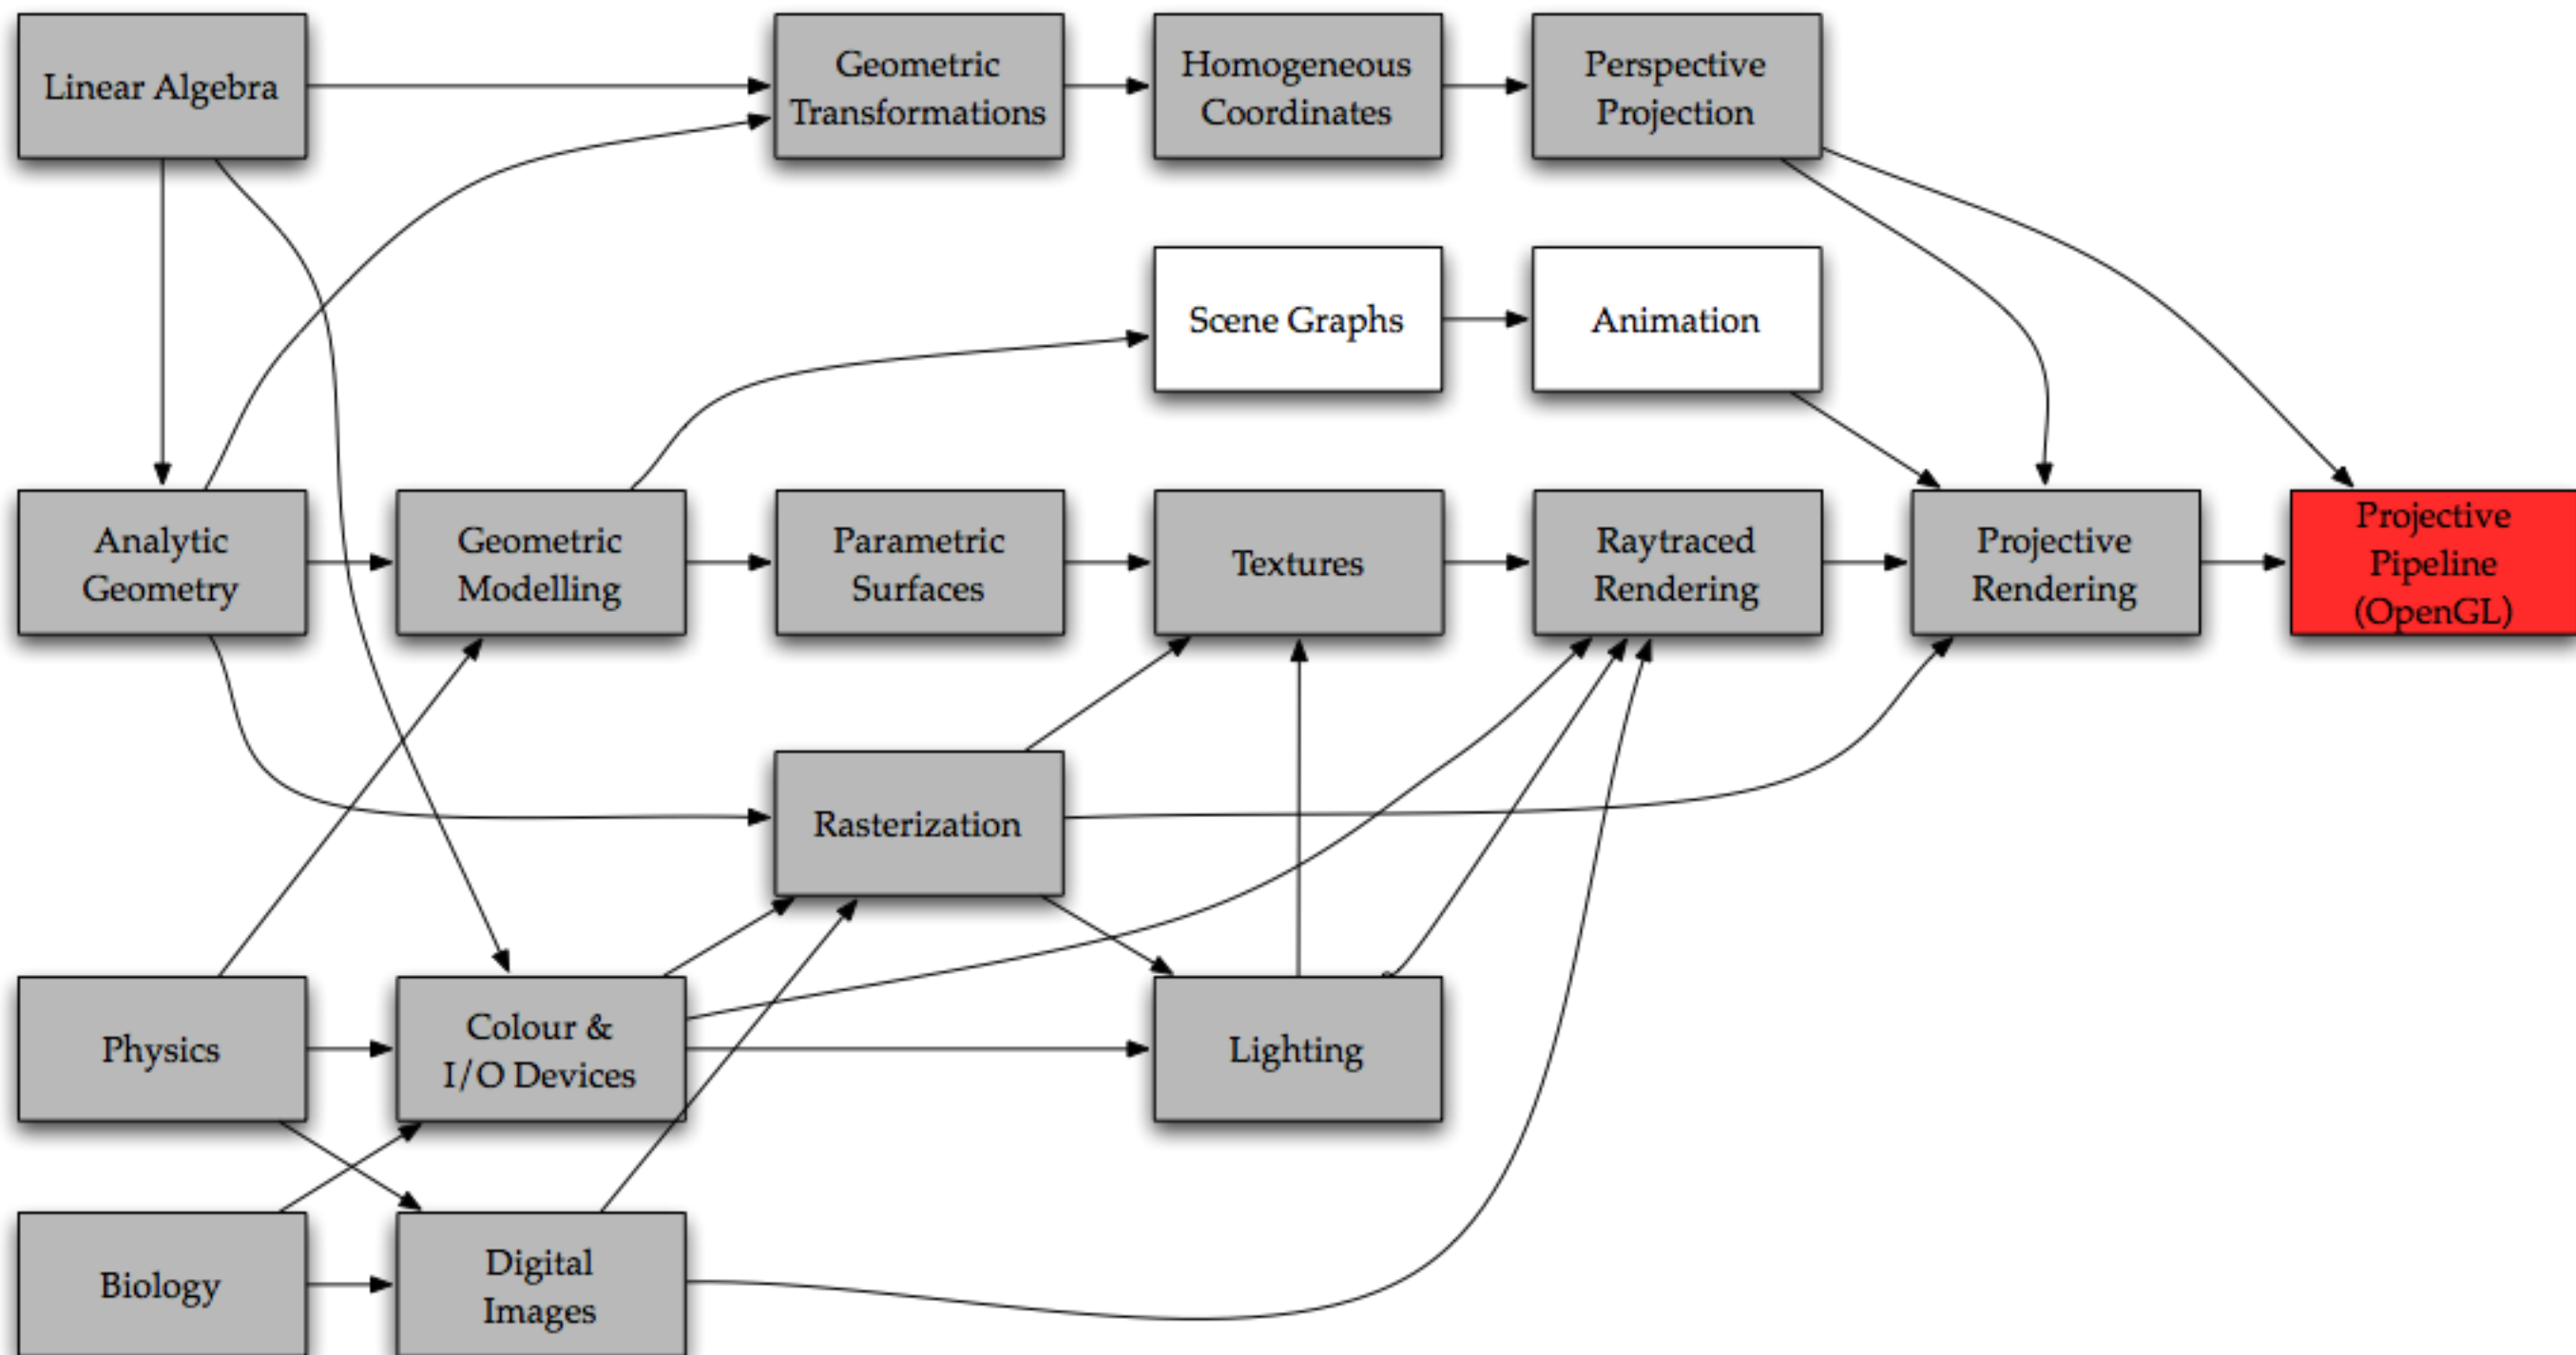


Coordinate Systems and Projective Rendering



Where we Are




Coordinate Systems

- OCS is the *Object Coordinate System*
- WCS is the *World Coordinate System*
- VCS is the *View Coordinate System*
- CCS is the *Clipping Coordinate System*
- NDCS is the *Normalized DCS*
- DCS is the *Device Coordinate System*



Object Coordinates

- Coordinates defining an object
 - e.g. the cylinder
 - travel with object



World Coordinates

- Arbitrary coordinate system
- Where is the origin?
 - The Earth?
 - The Sun?
 - Greenwich meridian?
- Used to keep track of other systems



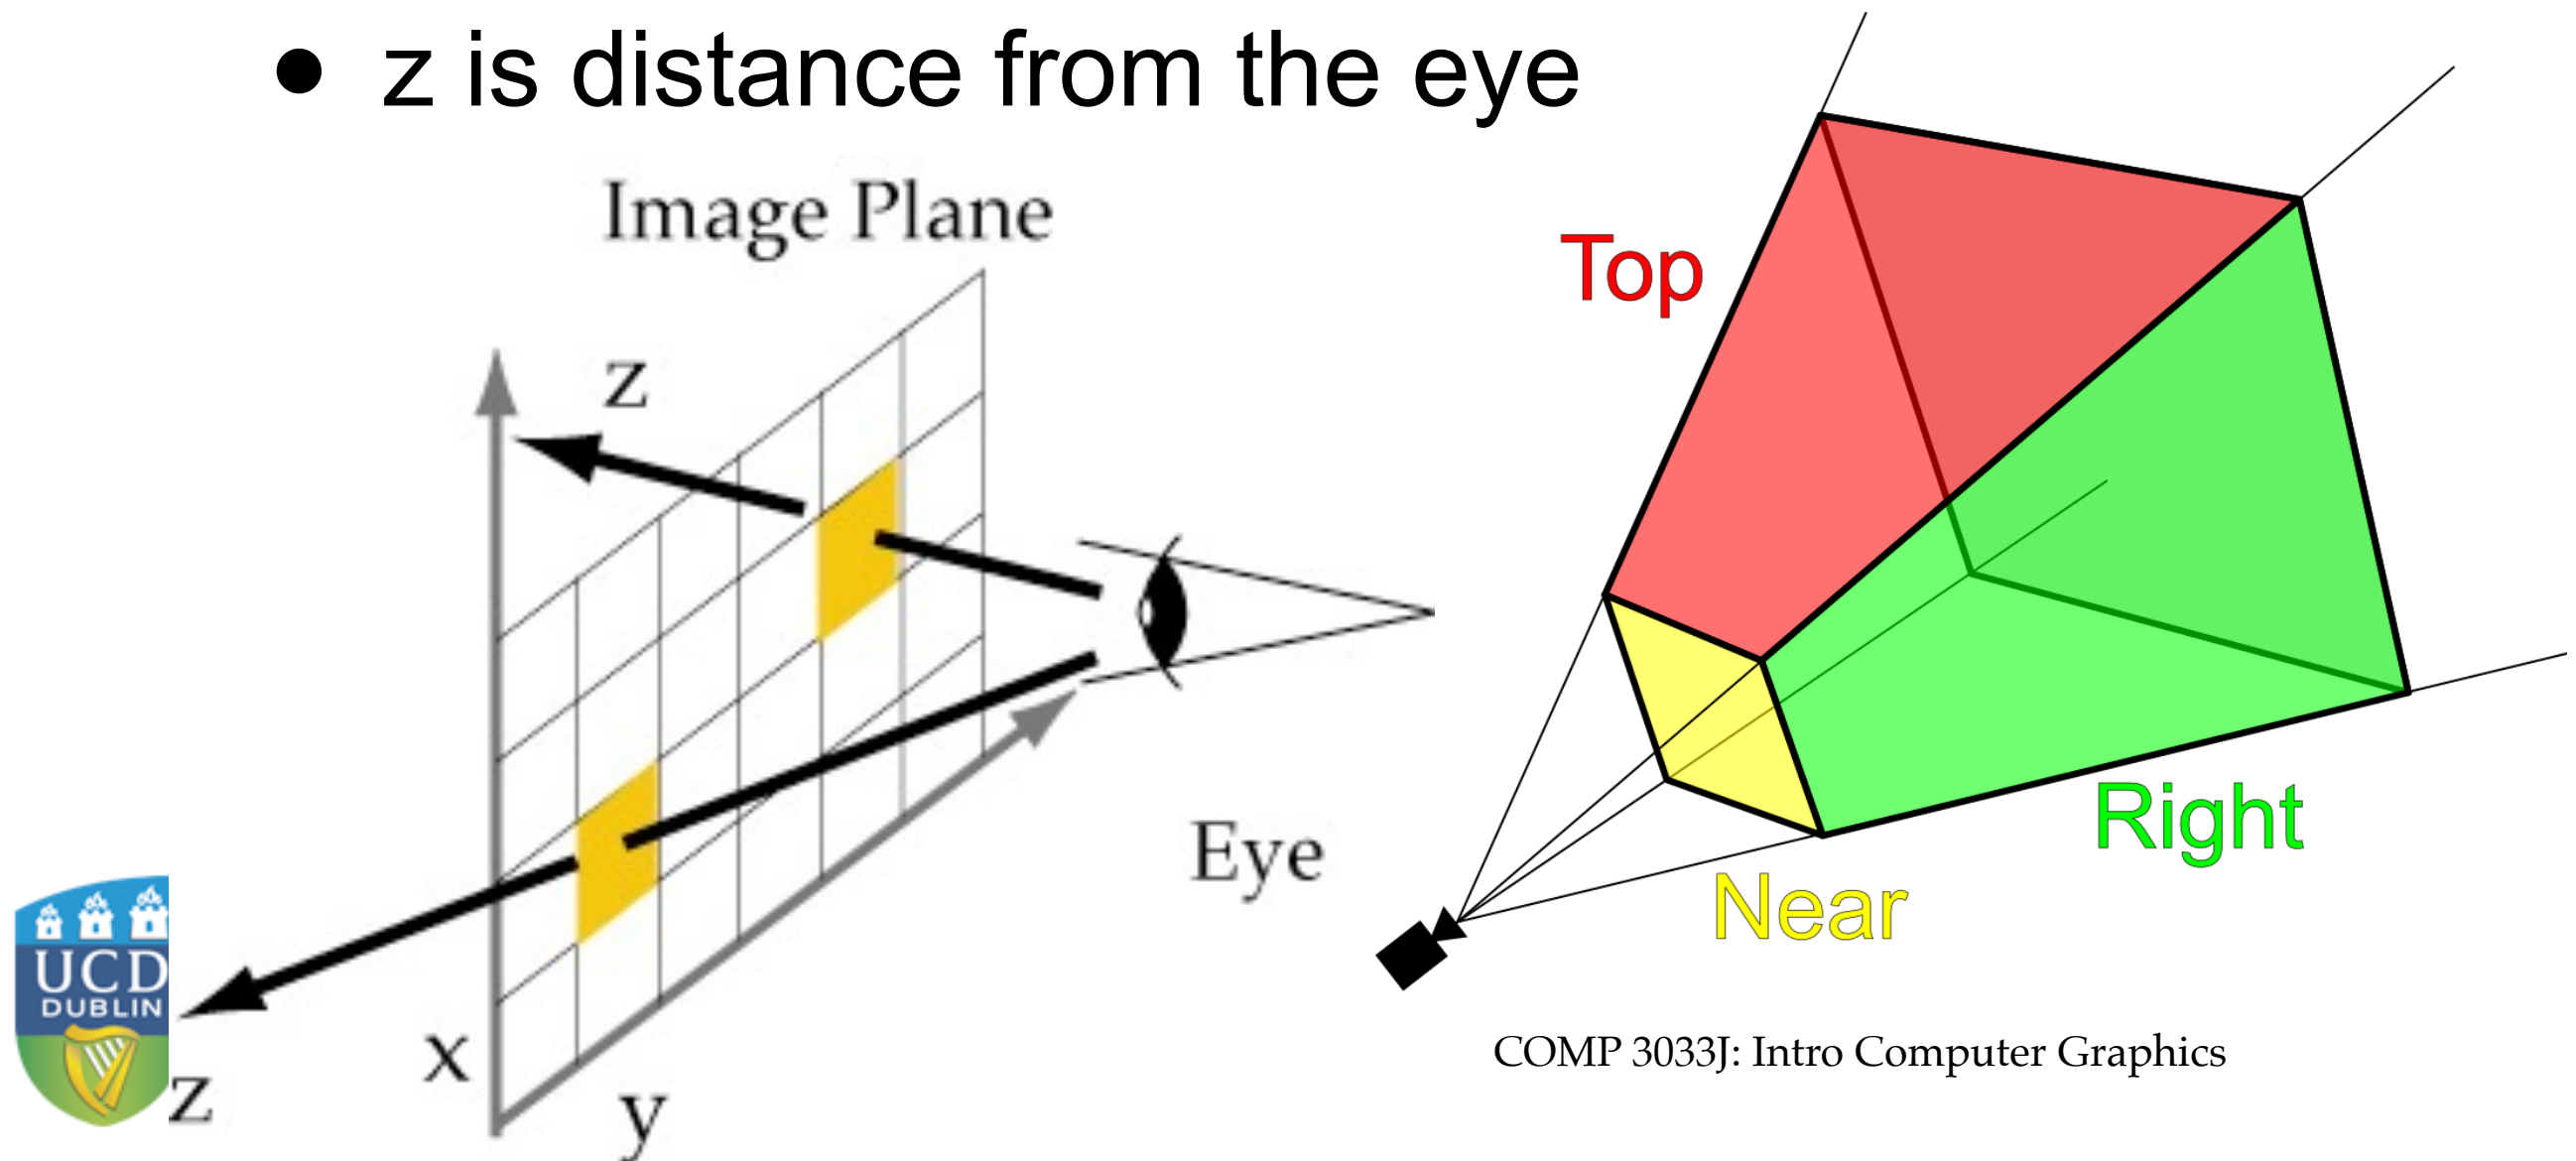
View Coordinates

- Tilt your head 90 degrees right
 - “Up” (y) is now to the right
 - “Backward” (z) hasn’t changed
 - We can use cross-product to get x
- Each camera / eye has coordinates



Clipping Coordinates

- coordinates in perspective projection
- z is distance from the eye



Clipping Coords

- Projection to CCS models the eye's lens
- In CCS, z is the distance from the eye
 - *not always in same direction*
 - i.e. Cartesian coordinates fail
 - homogeneous coordinates work
- View frustum is a box in these coords



Clipping Coords

Perspective Projection:

In perspective projection, the projection matrix is designed to simulate the way the human eye perceives depth. It maps points from 3D space to 2D space in such a way that objects farther from the camera appear smaller than those closer.



Clipping Coords

The general form of the perspective projection matrix looks like this:

$$P = \begin{pmatrix} \frac{1}{\tan(\frac{FOV_x}{2})} & 0 & 0 & 0 \\ 0 & \frac{1}{\tan(\frac{FOV_y}{2})} & 0 & 0 \\ 0 & 0 & \frac{z_f + z_n}{z_n - z_f} & \frac{2z_f z_n}{z_n - z_f} \\ 0 & 0 & -1 & 0 \end{pmatrix}$$

Where:

FOV_x, FOV_y represent the field of view angles in the horizontal and vertical directions.

z_n, z_f represent the near and far clipping planes, defining the depth range visible to the camera.



Normalized DCS

- *Normalized Device Coordinate System*
- Big mouthful, but simple idea
- Divide CCS through by w
 - converts homogeneous coords to Cartesian
- Independent of screen coordinates



Device Coordinates

- Finally, we convert into *device* coords
 - (x, y) : position on image plane
 - z : distance in front of image plane
- I.e. pixel position & object depth

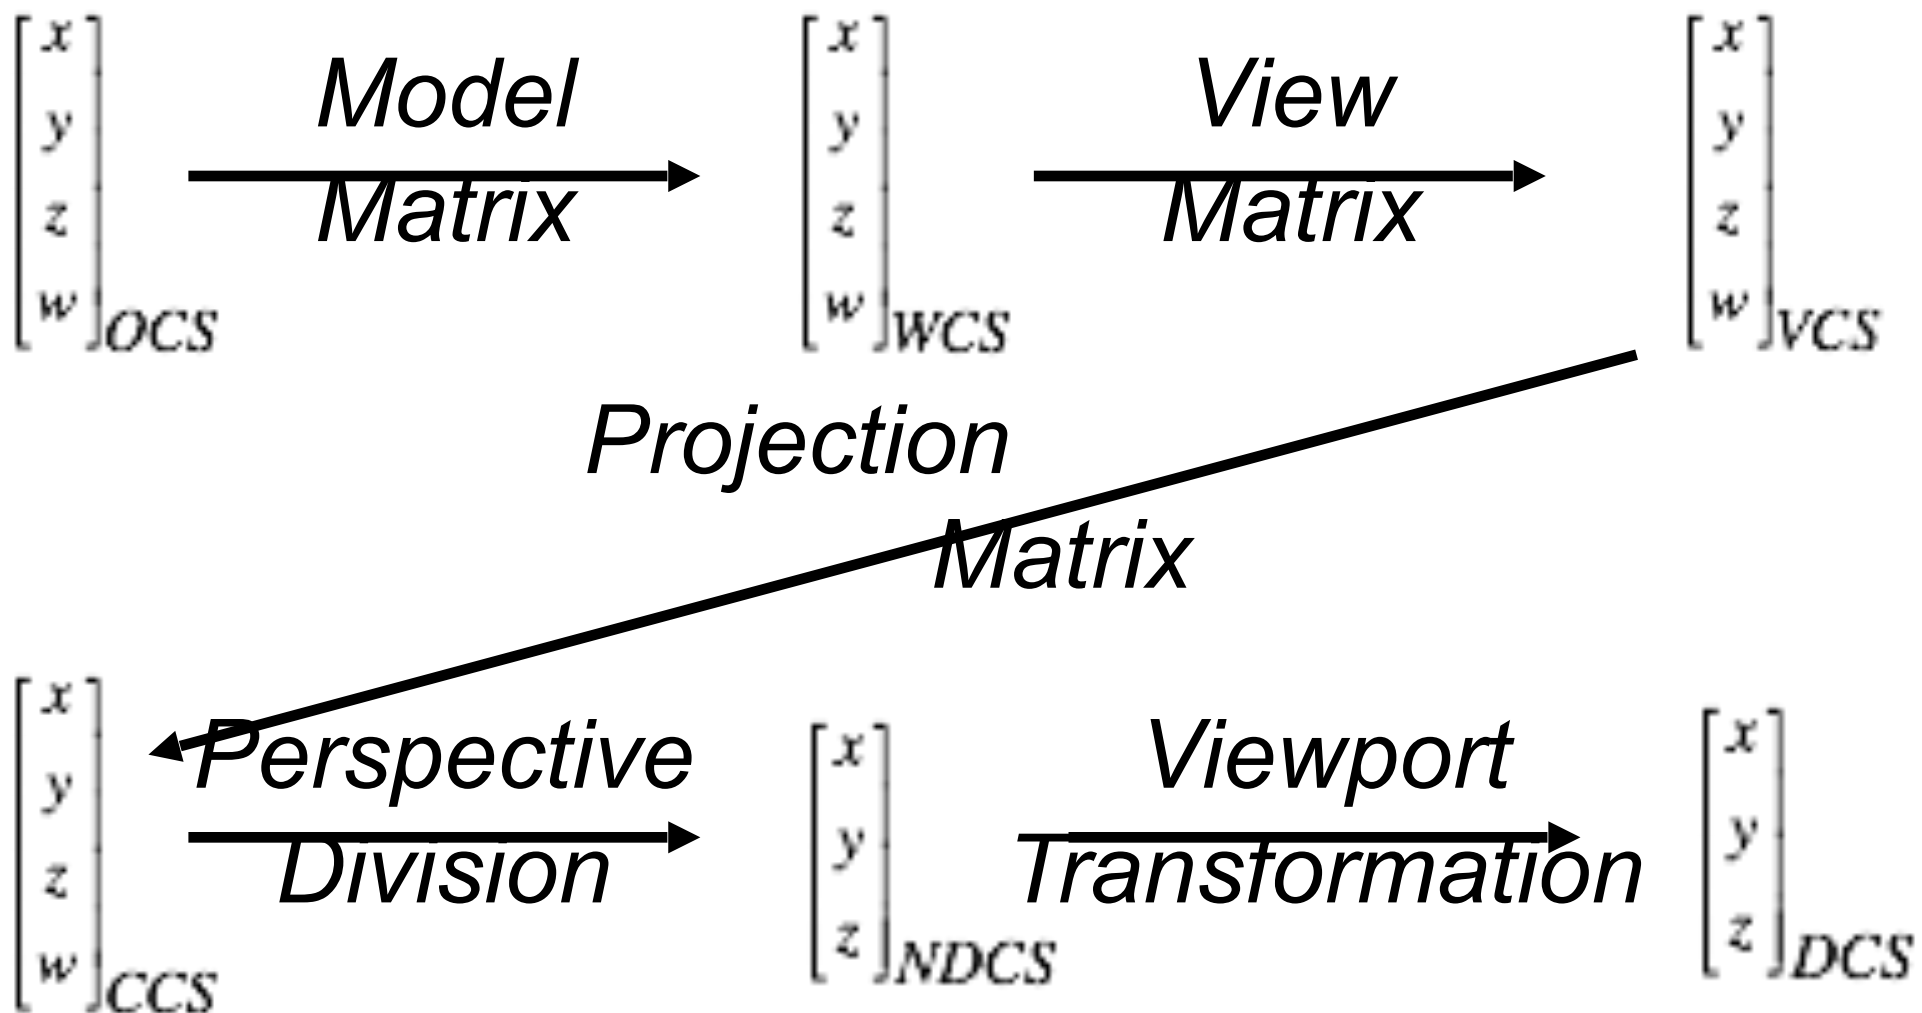


Why so many?

Object Model

World Model

View Model



Clipping

Image Plane

Screen

Transformations

- Model Matrix
- View Matrix
- Projection Matrix
- Perspective Division
- Viewport Transformation



Model Matrix

- Most commonly changed matrix
- Converts vertices from *OCS* to *WCS*
- *Or* moves *OCS* around in *WCS*
- Used to position & orient objects in scene



View Matrix

- Second most commonly changed matrix
- Converts from *WCS* to *VCS*
- *Or* moves *VCS* relative to *WCS*
- Gives position / orientation of *eye* in scene



Projection Matrix

- Less commonly changed
- Converts *VCS* to *CCS*
- Used to specify characteristics of *camera*
 - field of view, depth of field, symmetry
- Objects are usually *clipped* when in *CCS*
 - i.e. cut off at boundaries of frustum



Perspective Division

- Converts from homogeneous coordinates
- Gets rid of scale factor w
- Necessary for rasterization



Viewport Transformation

- Converts from *NDCS* to *DCS*
- *NDCS* is independent of window size
- *DCS* is not
- Used to separate rendering from O/S
- Change whenever window resized



Projective Rendering

- Raytracing is hard to accelerate
 - you don't know what the ray will hit
 - processing gets *complicated*
- Projective rendering is much easier
 - just throw *all* the objects at the screen
 - broken down into simpler steps



Basic Idea

- Each object is modelled geometrically
 - as *primitives* built from *vertices*
- Vertices are then *projected* to screen
 - i.e. *geometric* computations
- Primitives are *rasterized*
 - i.e. *pixel (fragment)* computations



Advantages

- Simple programming model
 - no complicated intersection tests
- Abstracts object definition
- Optimizes matrix multiplications
- Well-suited to parallelization
 - especially in dedicated hardware




Disadvantages

- Lots of setup required
- Opaque black-box computation
 - debugging is *hard*
 - but bugs are pretty standard
- Hard to keep track of *meaning* of matrices



Projective Rendering



I: Setup

- Define *model, view, projection* matrices
- Define *viewport* transformation
- Define *lights* & lighting properties
- Define *shading & lighting* model
- Load *textures* as images
- Set *texture* properties



II: Define Objects

- Describe object with *primitives*
 - points, lines, triangles
 - made up of *vertices*
- Define *colour, material properties, textures*
 - for each vertex



III: Geometric Operations

- *Transform* vertices to image coordinates
 - apply *model, view, projection* matrices
 - perform *perspective* division
 - apply *viewport* transformation
- Also transform *normal* vectors
- *Clip* primitives to view volume



IV: Rasterization

- Rasterize primitives to *fragments*
 - i.e. pixels which may or may not render
- Use barycentric coordinates to interpolate
 - colour
 - material properties
 - texture coordinates



V: Pixel Operations

- For each *fragment*
 - find *Texel* colour
 - with bilinear or NN interpolation
 - *combine* texture & shading
 - *replace* or *modulate*
 - perform *fragment* tests
 - Stencil (aids shadows) , depth buffer (is one object in front of another, &c.




VI: Framebuffer

- Write pixel to framebuffer
- Keep two framebuffers (at least)
 - for *double-buffering*
- Hardware transfers framebuffer to screen
 - or to an image
 - or to a texture



OpenGL 2 Pipeline



- From the OpenGL Red Book